

# Visualizing Data

Ismail Fadeli

August 19, 2021

## Data Visualization in R

In order to be able to visualize data, we have to install **tidyverse** package in R Studio. Here's the code for that:

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.3      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

## Viewing Dataset

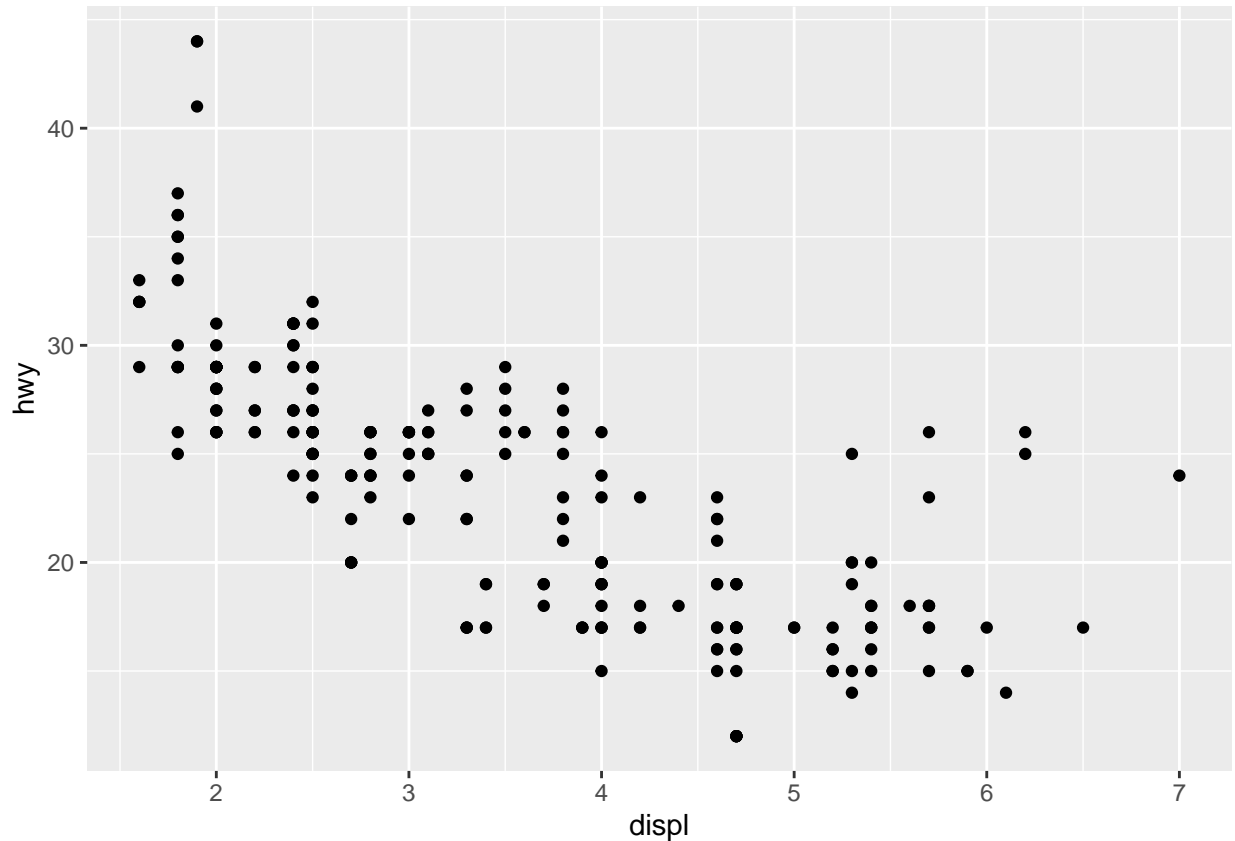
To visualize cars miles per gallon dataset from the USA datacenter:

```
mpg

## # A tibble: 234 x 11
##   manufacturer model      displ  year  cyl trans drv     cty   hwy fl      class
##   <chr>          <chr>    <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi          a4         1.8  1999    4 auto~ f      18    29 p    comp~
## 2 audi          a4         1.8  1999    4 manu~ f      21    29 p    comp~
## 3 audi          a4         2    2008    4 manu~ f      20    31 p    comp~
## 4 audi          a4         2    2008    4 auto~ f      21    30 p    comp~
## 5 audi          a4         2.8  1999    6 auto~ f      16    26 p    comp~
## 6 audi          a4         2.8  1999    6 manu~ f      18    26 p    comp~
## 7 audi          a4         3.1  2008    6 auto~ f      18    27 p    comp~
## 8 audi          a4 quattro 1.8  1999    4 manu~ 4      18    26 p    comp~
## 9 audi          a4 quattro 1.8  1999    4 auto~ 4      16    25 p    comp~
## 10 audi         a4 quattro 2    2008    4 manu~ 4      20    28 p    comp~
## # ... with 224 more rows
```

To plot **mpg** data, we need to run this code to put **displ** into x-axis and **hwy** into the y-axis.

```
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy))
```



The visualized graph shows a negative relationship between the engine size and fuel efficiency. Bigger engine size uses more fuel than lower engine sizes to travel the same distance.

With **ggplot**, you begin a plot with the function **ggplot()** creates a coordinate system that you can add layers to. The first argument of **ggplot** is the dataset to use in the graph. So **ggplot(data = mpg)** creates an empty graph. We can complete the graph by adding more layers to **ggplot()**. The function **geompoint()** creates a layer to your plot, which creates a scatterplot. The mapping argument is always paired with **aes()**, and the x and y arguments of **aes()** specify which variables to map to the x-axes and y-axes.

## Exercise

1. When running the code **ggplot(data = mpg)**, we see an empty graph:
2. The **mtcars** dataset has 32 rows and 11 columns.

**mtcars**

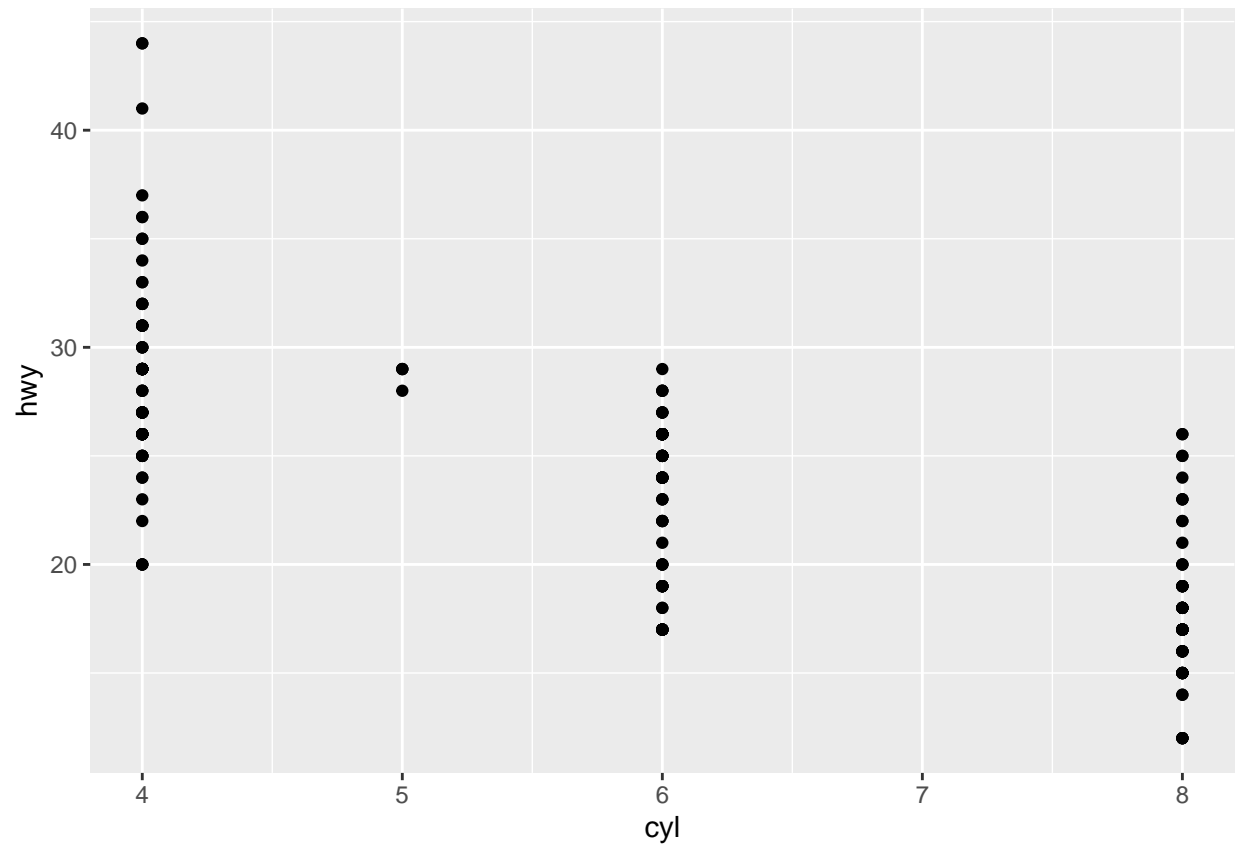
```
##           mpg  cyl  disp  hp  drat    wt   qsec  vs  am  gear  carb
## Mazda RX4      21.0    6  160.0  110  3.90  2.620  16.46  0   1    4     4
## Mazda RX4 Wag  21.0    6  160.0  110  3.90  2.875  17.02  0   1    4     4
## Datsun 710     22.8    4  108.0   93  3.85  2.320  18.61  1   1    4     1
## Hornet 4 Drive  21.4    6  258.0  110  3.08  3.215  19.44  1   0    3     1
## Hornet Sportabout 18.7    8  360.0  175  3.15  3.440  17.02  0   0    3     2
## Valiant        18.1    6  225.0  105  2.76  3.460  20.22  1   0    3     1
## Duster 360     14.3    8  360.0  245  3.21  3.570  15.84  0   0    3     4
## Merc 240D      24.4    4  146.7   62  3.69  3.190  20.00  1   0    4     2
## Merc 230       22.8    4  140.8   95  3.92  3.150  22.90  1   0    4     2
## Merc 280       19.2    6  167.6  123  3.92  3.440  18.30  1   0    4     4
## Merc 280C      17.8    6  167.6  123  3.92  3.440  18.90  1   0    4     4
```

## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

3. The **drv** variable describes the type of the car, meaning it is either a front wheel drive, rear wheel, or four wheel drive.

4. The code for a scatter plot of **hwy** versus **cyl**:

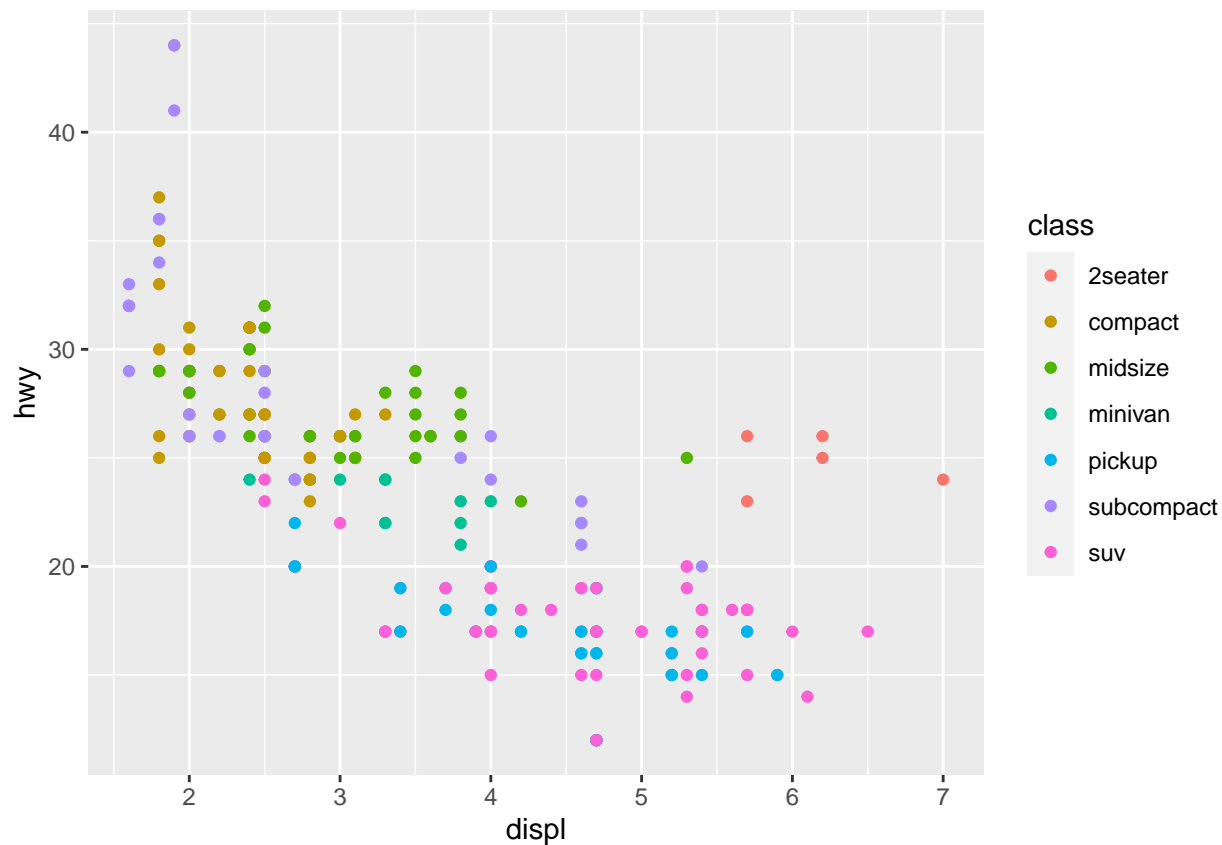
```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = cyl, y = hwy))
```



## Aesthetic Mappings

You can show information about your data by mapping the aesthetics in your plot to the variables in your dataset. For example, you can map the colors of your points to the class variable to reveal the class of each car:

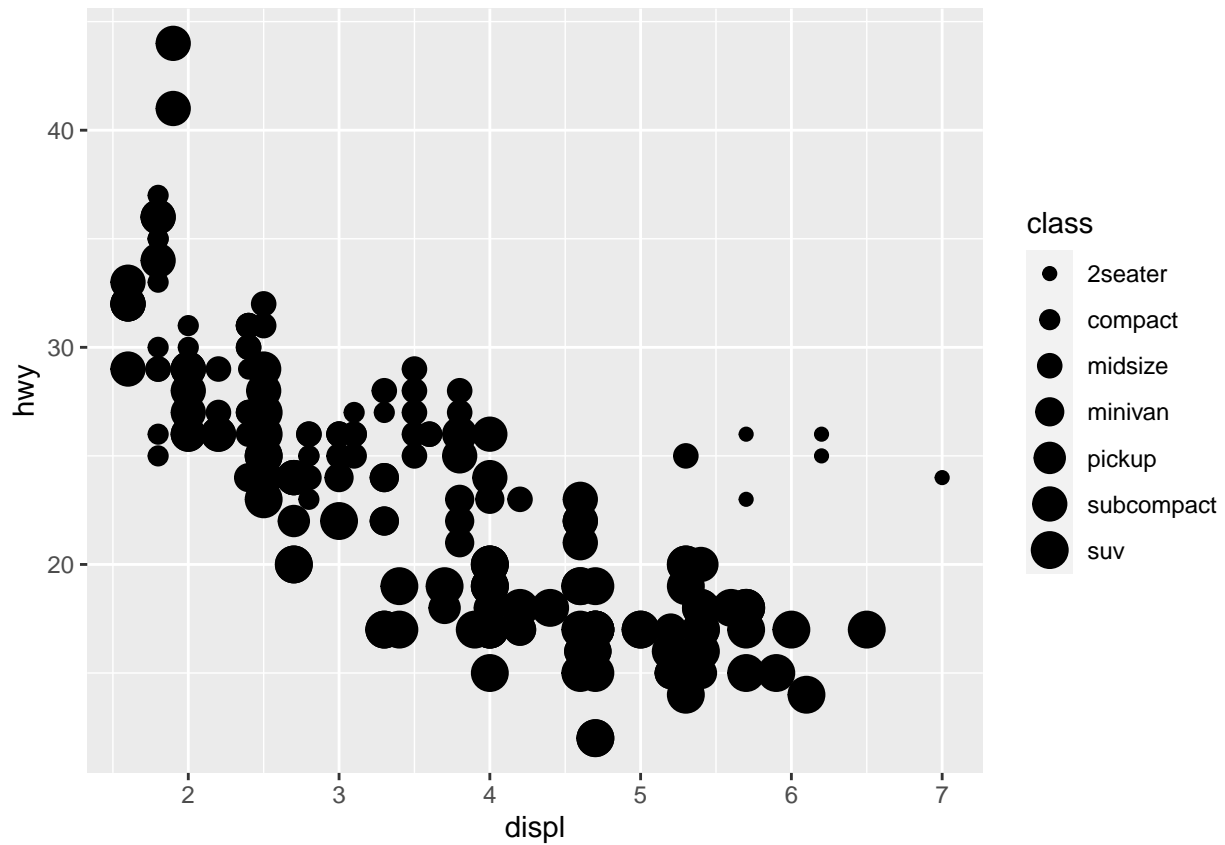
```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, color = class))
```



In the previous example, we mapped the class to the color aesthetic, but we could have mapped class to the size aesthetic in the same way. In this case, the exact size of each point would reveal its class affiliation. We get a warning here, because mapping an unordered variable (class) to an ordered variable (size) is not a good idea:

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, size = class))
```

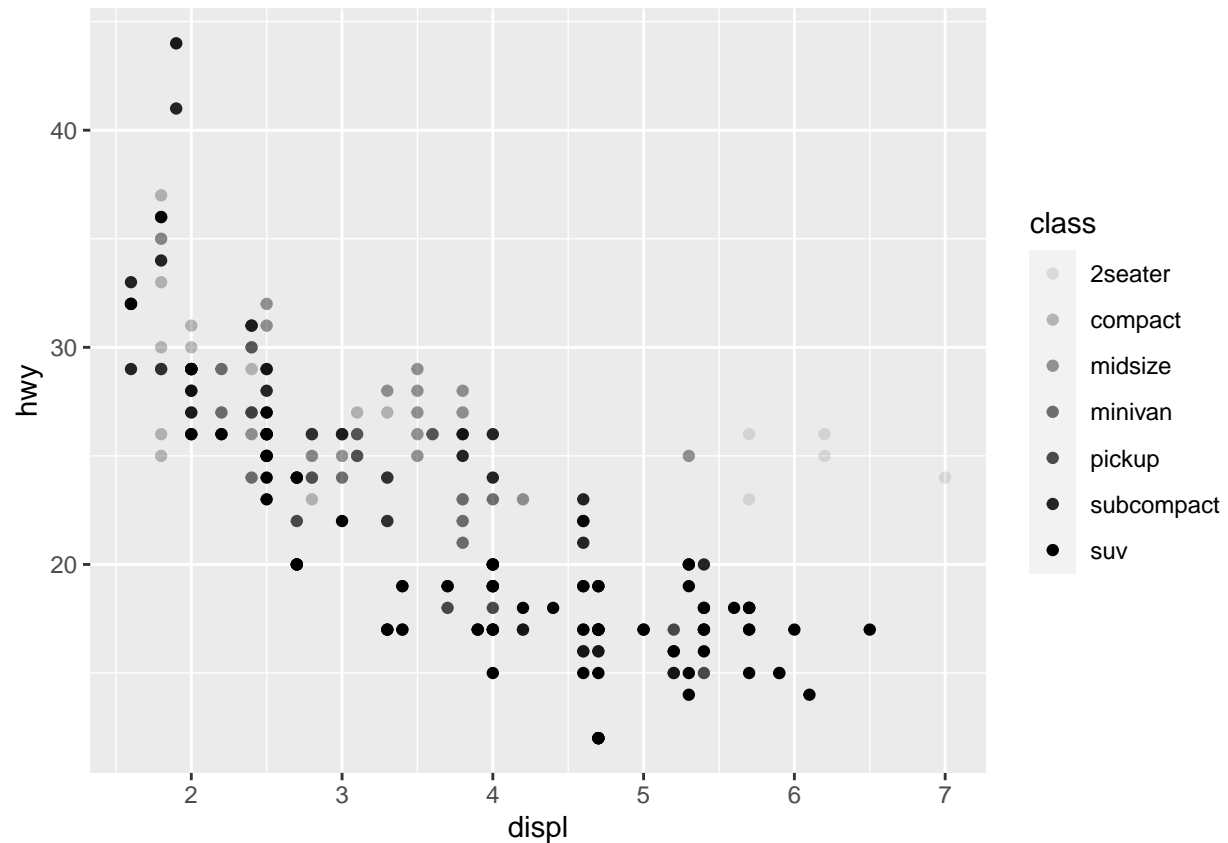
```
## Warning: Using size for a discrete variable is not advised.
```



Or we could have mapped class to the **alpha** aesthetic, which controls transparency of the points, or the shape of the points:


```
# Top
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy, alpha = class))
```

```
## Warning: Using alpha for a discrete variable is not advised.
```



```
# Bottom
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy, shape = class))
```

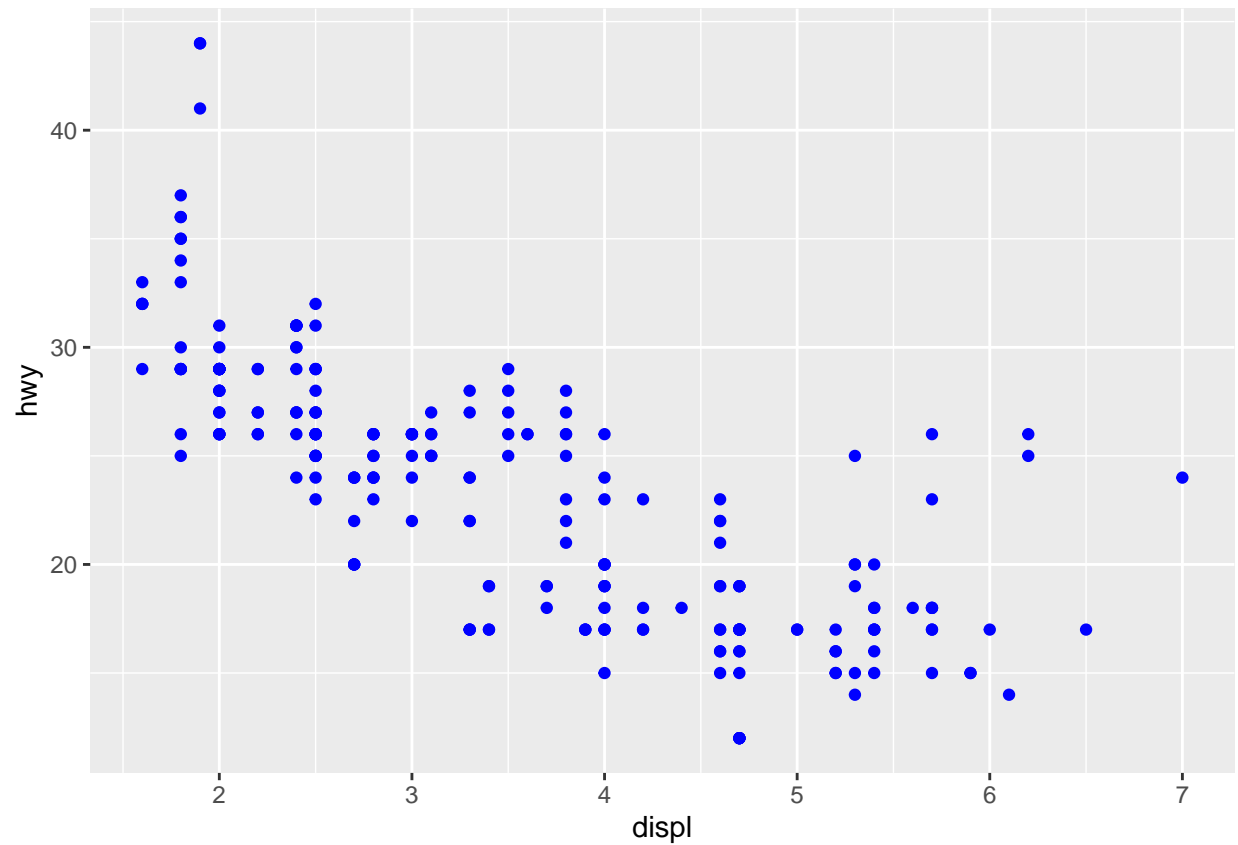
```
## Warning: The shape palette can deal with a maximum of 6 discrete values because
## more than 6 becomes difficult to discriminate; you have 7. Consider
## specifying shapes manually if you must have them.
## Warning: Removed 62 rows containing missing values (geom_point).
```



You can also set the aesthetic properties of your geom manually. For example, we can make all of the points in our plot blue:

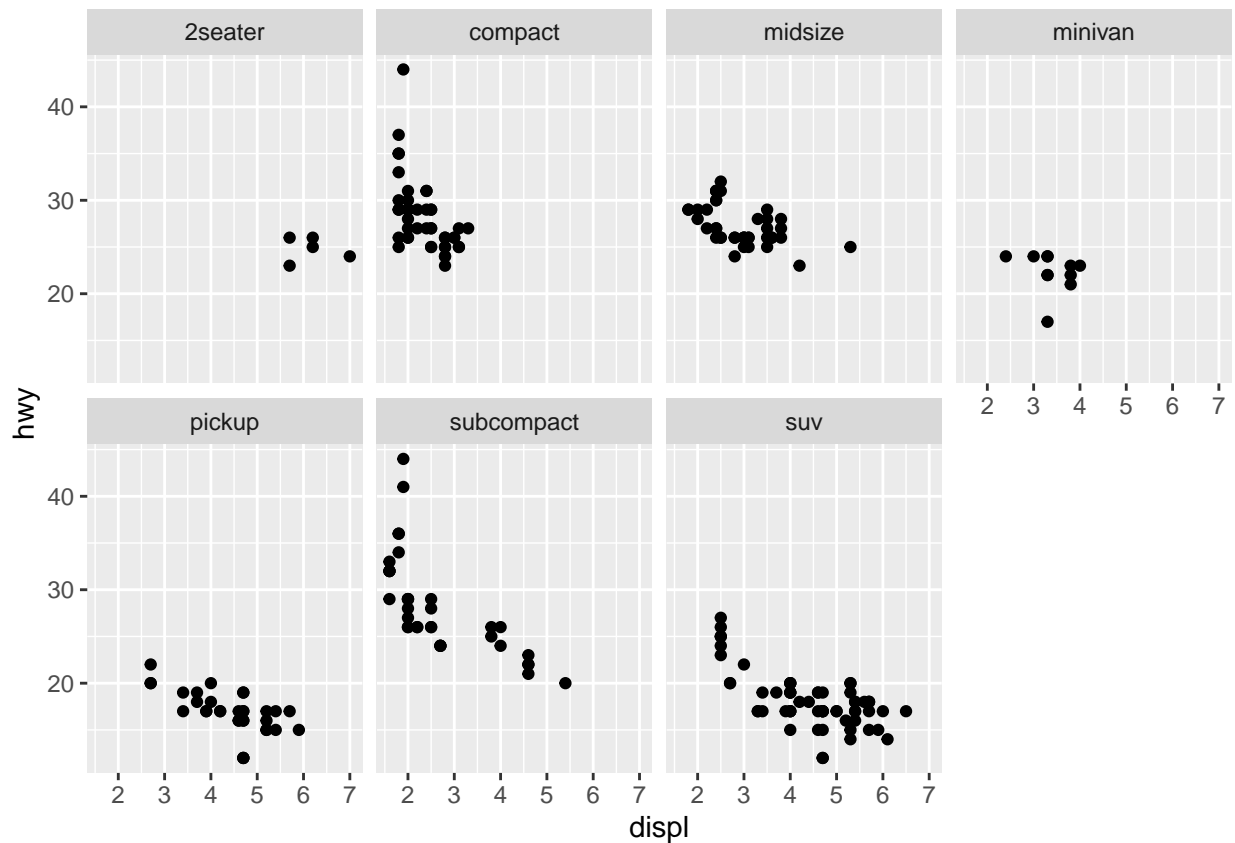
```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy), color = "blue")
```






## Facets one way to split categorical variables is to plot into facets, subplots that each display one subset of the data.

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy)) +  
  facet_wrap(~class, nrow = 2)
```



To facet the plot on combination of two variables, add **facetgrid** to your plot call. The first argument of **facetgrid** is also a formula. This time the formula should contain two variable names separated by a `~`:

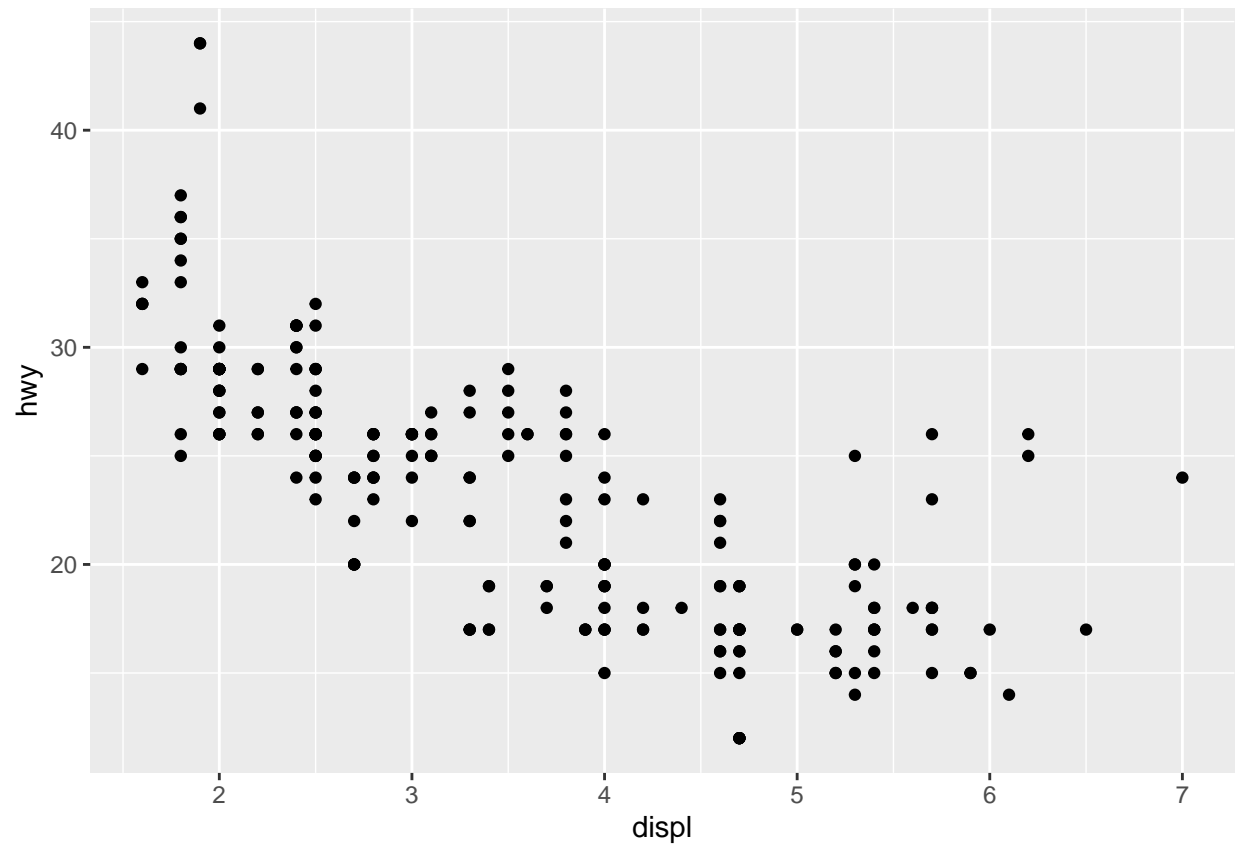
```
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy)) +
  facet_grid(drv ~ cyl)
```



## Geometric Objects


# Left

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))
```




```
# Right
ggplot(data = mpg) +
  geom_smooth(mapping = aes(x = displ, y = hwy))

## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, linetype = drv))
```


```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))  
  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, group = drv))  
  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = mpg) +  
  geom_smooth(  
    mapping = aes(x = displ, y = hwy, color = drv),  
    show.legend = FALSE  
  )
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```






```
ggplot(mpg) +  
  geom_text(mapping = aes(x = displ, y = hwy, label = model))
```



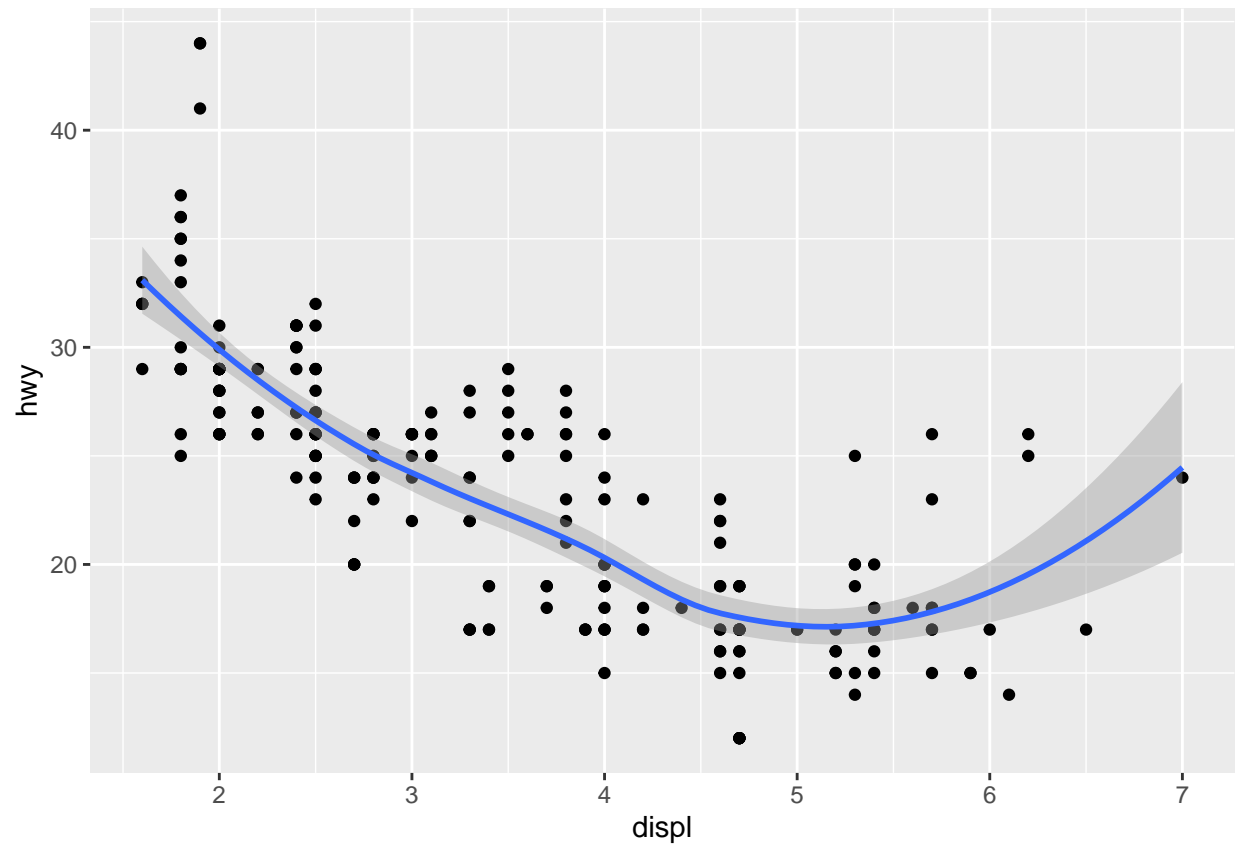
In order to display multiple geoms in the same plot, add multiple geom functions to `ggplot()`:

```
ggplot(data = mpg) +
  geom_point(mapping = aes(x = displ, y = hwy)) +
  geom_smooth(mapping = aes(x = displ, y = hwy))
```

## `geom\_smooth()` using method = 'loess' and formula 'y ~ x'

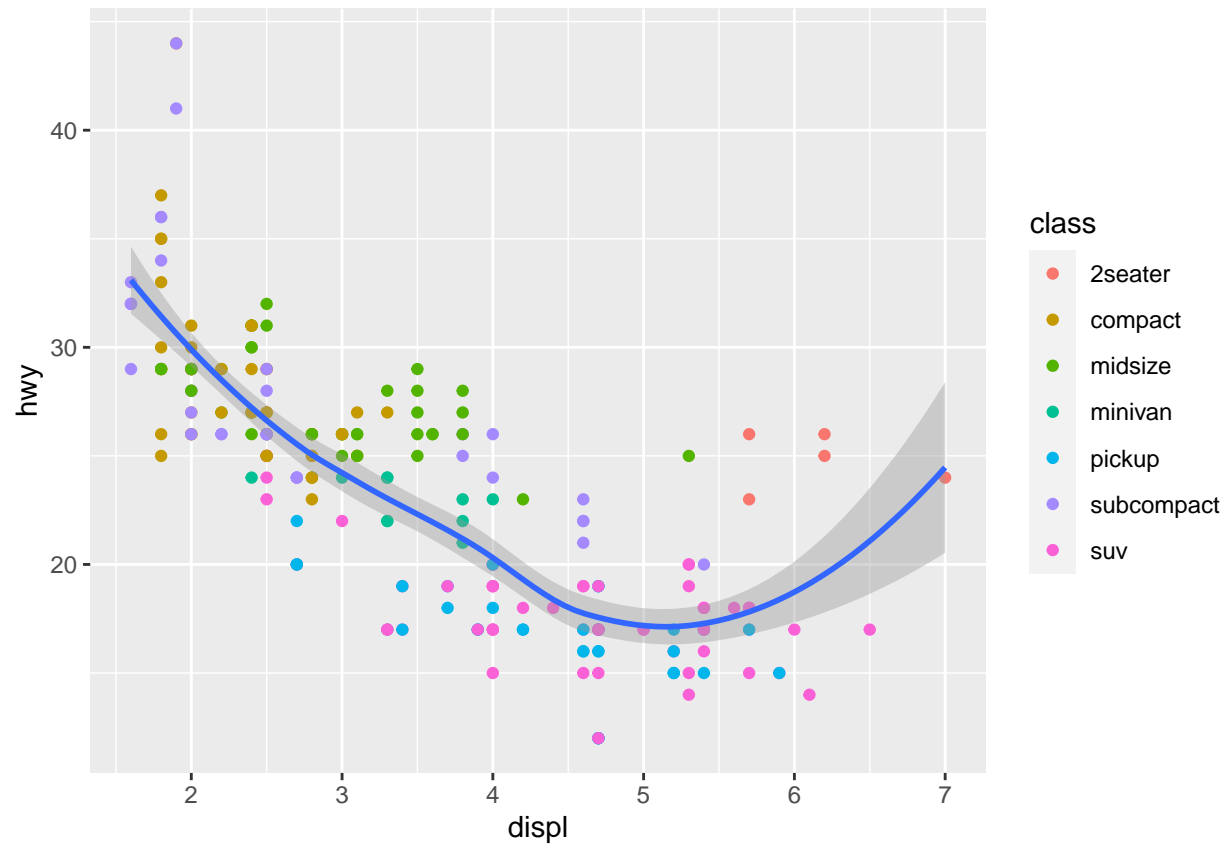


```
ggplot(data = mpg, mapping = aes( x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()  
  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(data = mpg, mapping = aes( x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth()
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth(  
    data = filter(mpg, class == "subcompact"),  
    se = FALSE  
  )  
  
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
diamonds
```

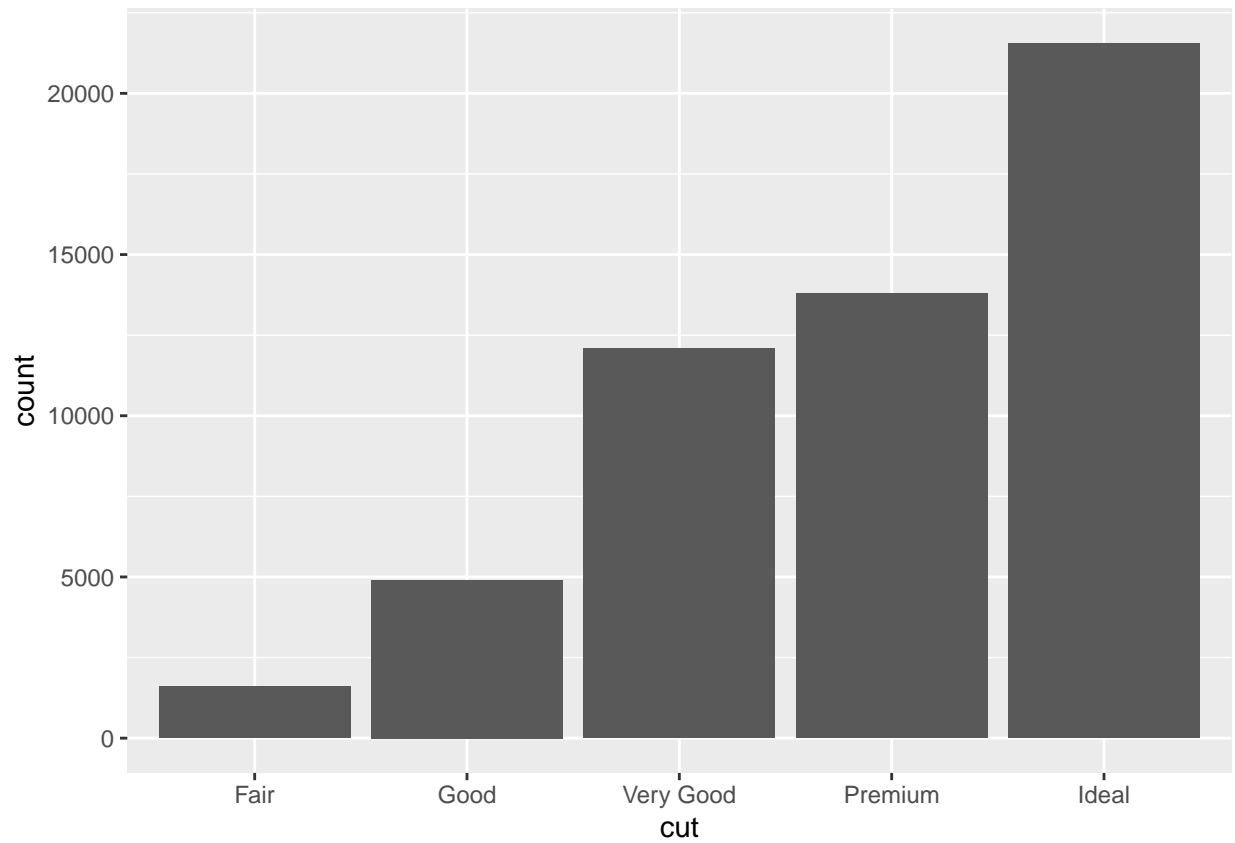
```
## # A tibble: 53,940 x 10
##   carat cut      color clarity depth table price      x      y      z
##   <dbl> <ord>    <ord> <ord>  <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1 0.23 Ideal    E     SI2    61.5   55   326  3.95  3.98  2.43
## 2 0.21 Premium E     SI1    59.8   61   326  3.89  3.84  2.31
## 3 0.23 Good    E     VS1    56.9   65   327  4.05  4.07  2.31
## 4 0.29 Premium I     VS2    62.4   58   334  4.2   4.23  2.63
## 5 0.31 Good    J     SI2    63.3   58   335  4.34  4.35  2.75
## 6 0.24 Very Good J     VVS2    62.8   57   336  3.94  3.96  2.48
## 7 0.24 Very Good I     VVS1    62.3   57   336  3.95  3.98  2.47
## 8 0.26 Very Good H     SI1    61.9   55   337  4.07  4.11  2.53
## 9 0.22 Fair    E     VS2    65.1   61   337  3.87  3.78  2.49
## 10 0.23 Very Good H     VS1    59.4   61   338  4     4.05  2.39
## # ... with 53,930 more rows
```

## Statistical Transformation

```
ggplot(data = diamonds) +
  geom_bar(mapping = aes(x = cut))
```




```
ggplot(data = diamonds) +  
  stat_count(mapping = aes(x = cut))
```



You might want to override the default mapping from transformed variables to aesthetics. For example, you might want to display a bar chart of proportion, rather than count:

```
ggplot(data = diamonds) +  
  geom_bar(  
    mapping = aes(x = cut, y = ..prop.., group = 1)  
  )
```





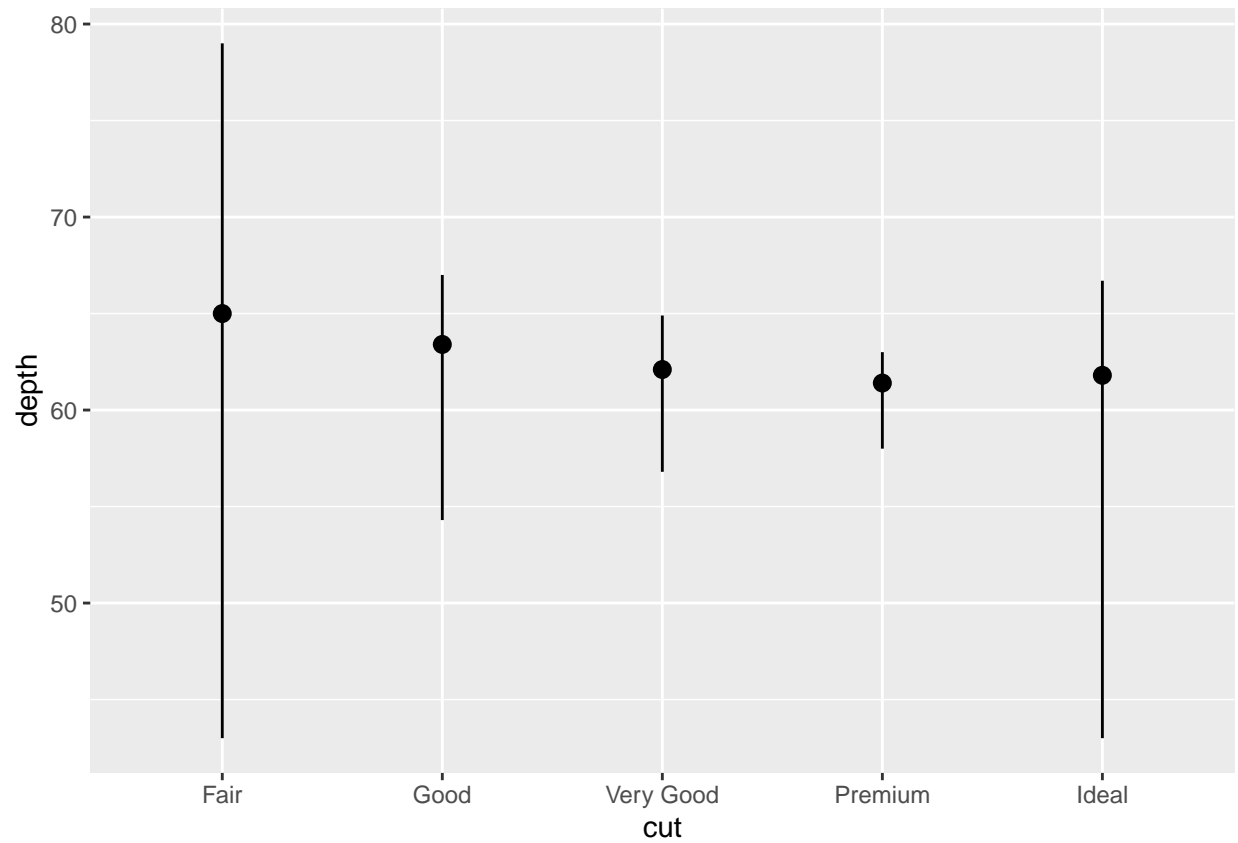
You might want to draw greater attention to the statistical transformation in your code. For example, you might use `stat_summary()`, which summarizes the y values for each unique x value, to draw attention to the summary that you're computing

```
ggplot(data = diamonds) +  
  stat_summary(  
    mapping = aes(x = cut, y = depth),  
    fun.ymin = min,  
    fun.ymax = max,  
    fun.y = median  
  )
```

```
## Warning: `fun.y` is deprecated. Use `fun` instead.
```


```
## Warning: `fun.ymin` is deprecated. Use `fun.min` instead.
```

```
## Warning: `fun.ymax` is deprecated. Use `fun.max` instead.
```




##Position Adjustments You can color a bar chart using either the color aesthetic, or more usefully, fill:

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, color = cut))
```




```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = cut))
```




Note what happens if you map the fill aesthetic to another variable, like clarity: the bars are automatically stacked. Each colored rectangle represents a combination of cut and clarity:

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity))
```



position = "identity" will place each object exactly where it falls in the context of the graph. This is not very useful for bars, because it overlaps them. To see that overlapping we either need to make the bars slightly transparent by setting alpha to a small value, or completely transparent by setting fill = NA:

```
ggplot(
  data = diamonds,
  mapping = aes(x = cut, fill = clarity)
) +
  geom_bar(alpha = 1/5, position = "identity")
```




```
ggplot(  
  data = diamonds,  
  mapping = aes(x = cut, color = clarity)  
)
```



position = “fill” works like stacking, but makes each set of stacked bars the same height. This makes it easier to compare proportions across groups:


```
ggplot(data = diamonds) +  
  geom_bar(  
    mapping = aes(x = cut, fill = clarity),  
    position = "fill"  
  )
```



position = "dodge" places overlapping objects directly beside one another. This makes it easier to compare individual values:

```
ggplot(data = diamonds) +
  geom_bar(
    mapping = aes(x = cut, fill = clarity),
    position = "dodge"
  )
```






here's one other type of adjustment that's not useful for bar charts, but it can be very useful for scatterplots. Recall our first scatterplot. Did you notice that the plot displays only 126 points, even though there are 234 observations in the dataset?


The values of `hwy` and `displ` are rounded so the points appear on a grid and many points overlap each other. This problem is known as overplotting. This arrangement makes it hard to see where the mass of the data is. Are the data points spread equally throughout the graph, or is there one special combination of `hwy` and `displ` that contains 109 values

You can avoid this gridding by setting the position adjustment to "jitter." `position = "jitter"` adds a small amount of random noise to each point. This spreads the points out because no two points are likely to receive the same amount of random noise:


```
ggplot(data = mpg) +
  geom_point(
    mapping = aes(x = displ, y = hwy),
    position = "jitter"
  )
```



```
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +  
  geom_boxplot()
```




```
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +  
  geom_boxplot() +  
  coord_flip()
```




```
nz <- map_data("nz")

ggplot(nz, aes(long, lat, group = group)) +
  geom_polygon(fill = "white", color = "black")
```



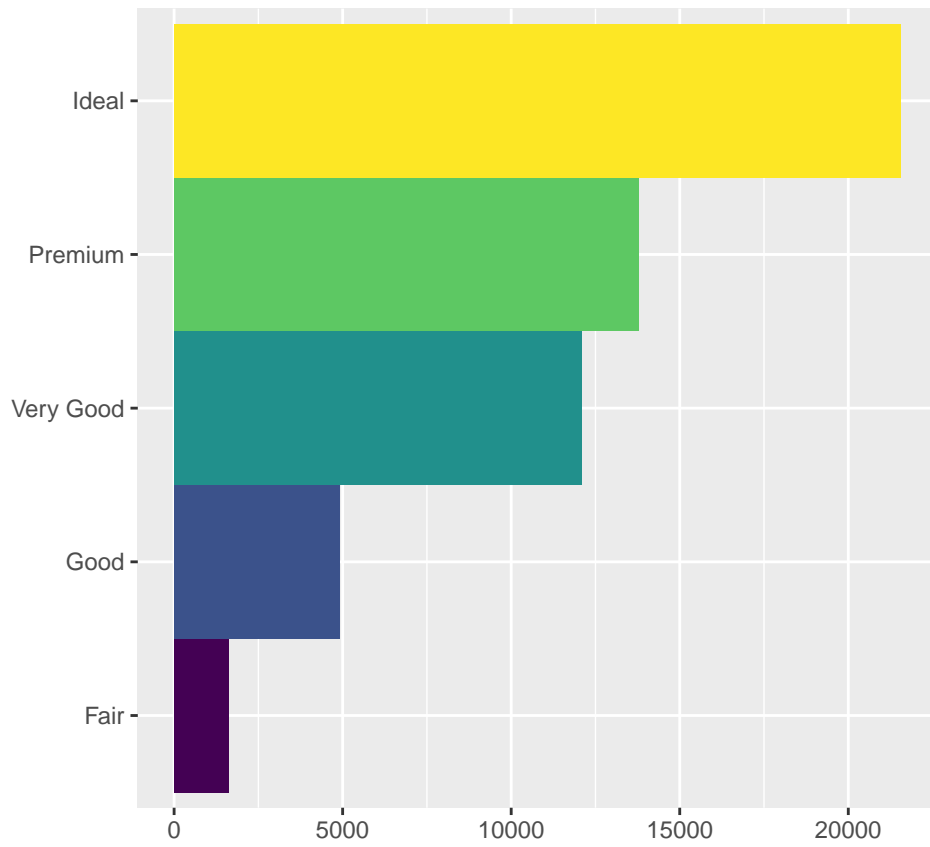
`coord_quickmap()` sets the aspect ratio correctly for maps. This is very important if you're plotting spatial data with `ggplot2` (which unfortunately we don't have the space to cover in this book):

```
ggplot(nz, aes(long, lat, group = group)) +  
  geom_polygon(fill = "white", color = "black") +  
  coord_quickmap()
```




`coord_polar()` uses polar coordinates. Polar coordinates reveal an interesting connection between a bar chart and a Coxcomb chart:

```
bar <- ggplot(data = diamonds) +  
  geom_bar(  
    mapping = aes(x = cut, fill = cut),  
    show.legend = FALSE,  
    width = 1  
  ) +  
  theme(aspect.ratio = 1) +  
  labs(x = NULL, y = NULL)  
  
bar + coord_flip()
```



```
bar + coord_polar()
```



```
my_val <- 10  
my_val
```

```
## [1] 10
```